

Parallel Computing for Science & Engineering

1/29/08

Instructors:

Dr. Bill Barth, Research Associate, TACC

Dr. Kent Milfeld, Research Associate, TACC

Programming with OpenMP on Shared Memory Systems

Reminders

- Enhanced performance and solving larger problems are the main reasons to use parallel computers
- Application developers have to design and program correct and efficient parallel codes for parallel computers to realize these benefits
- Achieving good single-processor performance is hard; achieving scalable parallel performance is harder
- But the payoff can be huge

OpenMP – what is it

- **De facto open standard for Scientific Parallel Programming on Symmetric MultiProcessor (SMP) Systems.**
- **Extensions to Fortran, C and C++ for portable SMP programming**
 - Bases on threads, but
 - Higher-level than POSIX threads (Pthreads)
(<http://www.llnl.gov/computing/tutorials/pthreads/#Abstract>)
- **Implemented by:**
 - **Compiler Directives**
 - **Runtime Library (an API, Application Program Interface)**
 - **Environment Variables**
- **Compiler option required to detect code directives.**
- **<http://www.openmp.org/> has tutorials and description.**

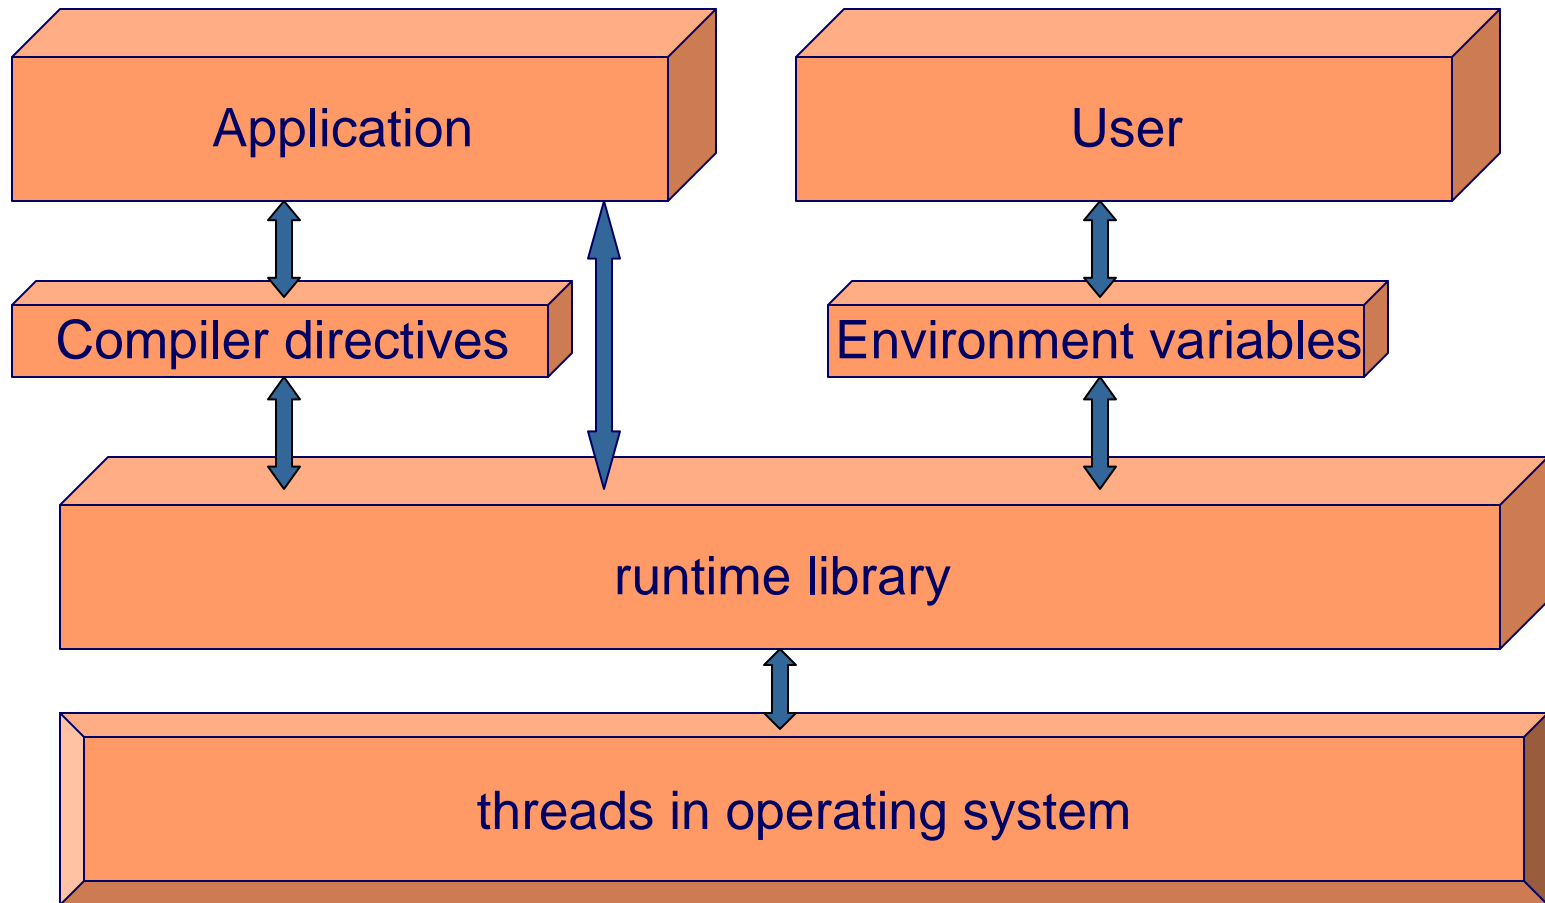
OpenMP History

- Primary OpenMP participants
Compaq, HP, IBM, Intel, KAI, SGI, SUN
- U.S. Dept. of Energy ASCI Program
- OpenMP Fortran API, Version 1.0, published Oct. 1997
- OpenMP C API, Version 1.0, published Oct. 1998
- OpenMP 2.0 API for Fortran, published in 2000
- OpenMP 2.0 API for C/C++, published in 2002
- OpenMP 2.5 API for C/C++ & F90 published in 2005

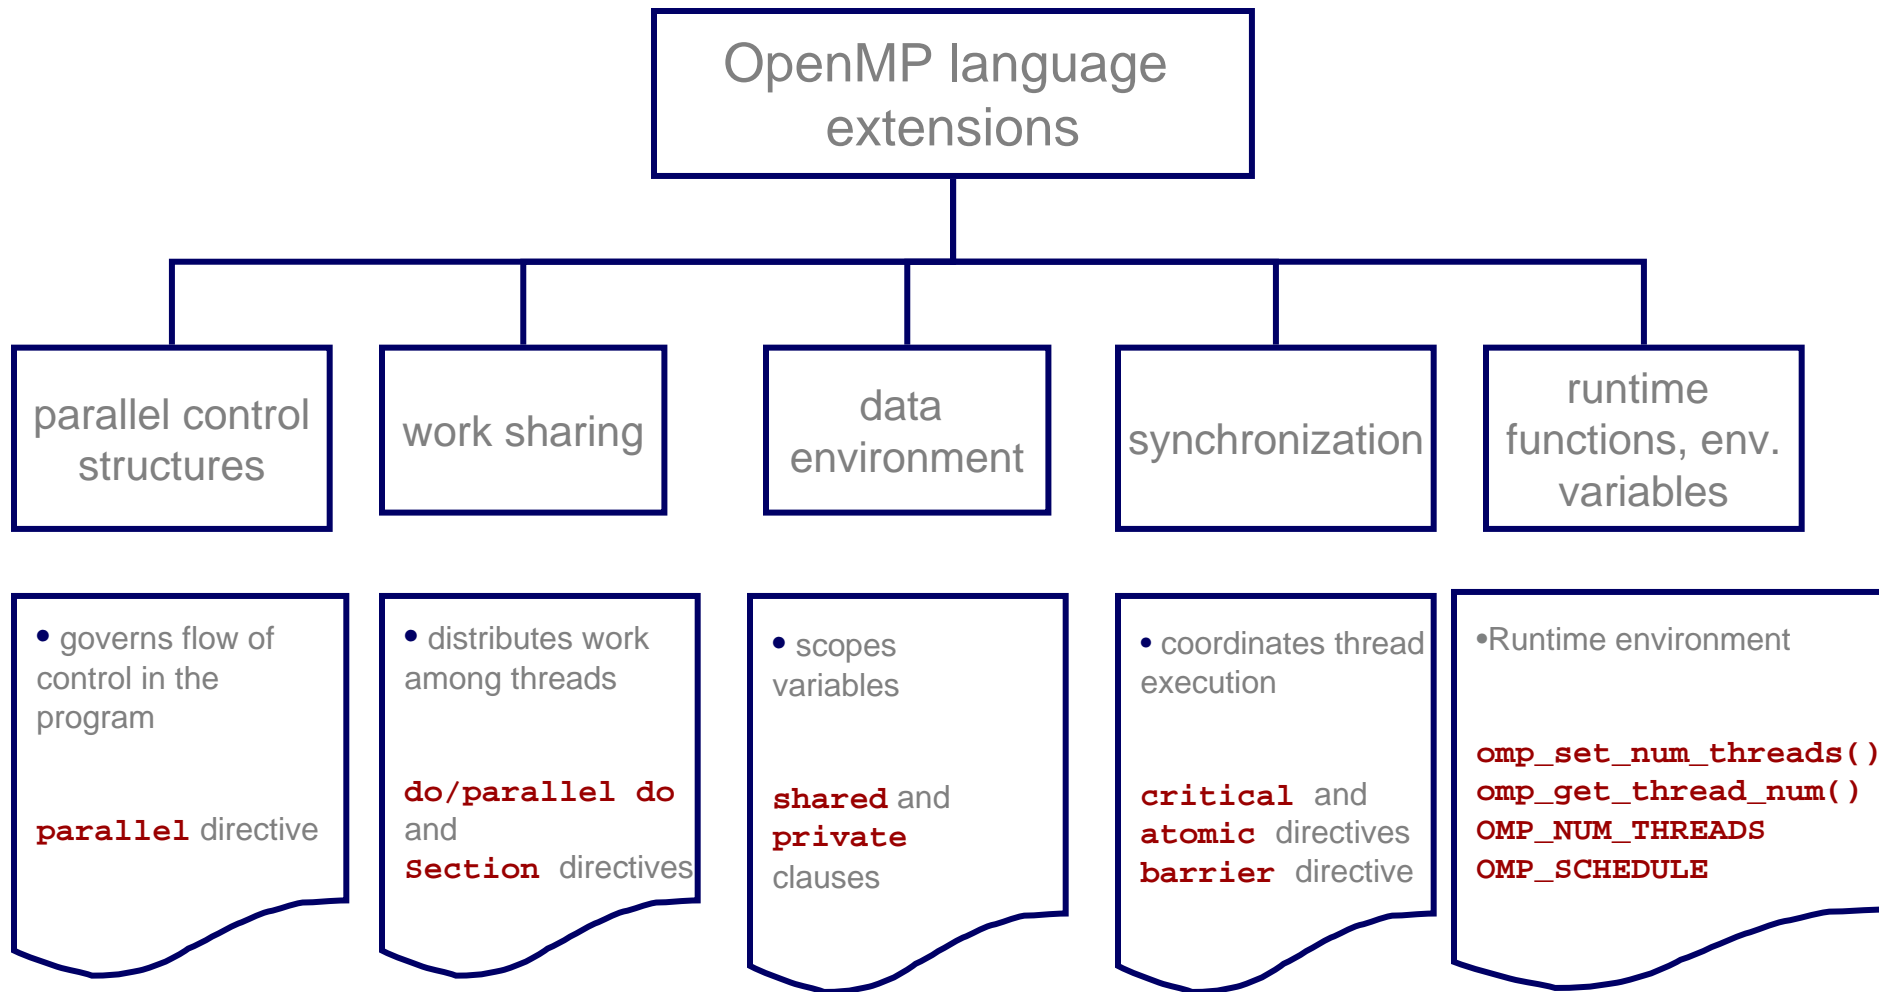
Advantages/Disadvantages of OpenMP

- Pros
 - Shared Memory Parallelism is easier to learn.
 - Coarse-grained or fine-grained parallelism
 - Parallelization can be incremental
 - Widely available, portable
- Cons
 - Scalability limited by memory architecture
 - Available on SMP systems only

OpenMP Architecture



OpenMP Constructs



OpenMP Parallel Directive

Fortran directives begin with: **!\$OMP**, **C\$OMP** or ***\$OMP** sentinel.

C/C++ directives begin with: **# pragma omp** sentinel.

Fortran Parallel regions are enclosed by **enclosing directives**

C/C++ Parallel regions are enclosed by **curly brackets**.

Syntax: *sentinel parallel clauses*

uses defaults w.o. clauses

Fortran

```
!$OMP parallel
...
!$OMP end parallel
```

C/C++

```
# pragma omp parallel
{...}
```


Parallel Region

1	!\$OMP PARALLEL	#pragma omp
2	<i>code block</i>	{ <i>code block</i>
3	call work(...)	i=work(...)
4	!\$OMP END PARALLEL	}

- Line 1** Team of threads formed at parallel region.
- Lines 2-3** Each thread executes code block and subroutine calls.
No branching (in or out) in a parallel region.
- Line 4** All threads synchronize at end of parallel region
(implied barrier).

OpenMP fork-join parallelism

- Parallel Regions are basic “blocks” within code.
- A master thread is instantiated at run-time & persists throughout execution.
- Master thread assembles team of threads at parallel regions.



OpenMP Clauses

Clauses control the behavior of directives

Syntax: *sentinel* parallel **clauses** *uses defaults w.o. clauses*

Data Scoping

private(list), shared(list), default(private|shared|none)

Scheduling

schedule(type,chunk)

IF Control

if(expression)

Ordering

ordered

Initialization

copyin (list), firstprivate(list), lastprivate(list)

Reduction

reduction (operator |intrinsic : list)

Barrier


nowait

How do threads communicate?

- Every thread has access to “global” memory (**shared**).
- Each thread has access to a stack memory (**private**).
- Use shared memory to communicate between threads.
- Simultaneous updates to shared memory can create a *race condition*. Results change with different thread scheduling.
- Use mutual exclusion to avoid data sharing --- but don't use too many because this will “serialize” performance.

Constructs within Parallel Regions

- Use OpenMP directives to specify Parallel Region and Work-Sharing constructs.



Parallel `DO/for`
Parallel `SECTIONS`

Stand-alone
Parallel Constructs

OpenMP Combined Directives

- Combined directives
 - PARALLEL DO/FOR and PARALLEL SECTIONS
 - Same as PARALLEL region containing only DO or SECTIONS work-sharing construct

```
!$omp parallel do  
  do i = 1, 100  
    a(i) = b(i)  
  end do
```

```
#pragma omp for  
for(i=0;i<100;i++){  
  a[i] = b[i];  
}
```

Work Sharing – do/for

Work-sharing (WS) constructs: do/for, section, and single

- WS Threads execution their “share” of statements in a PARALLEL region.
- Do/for Work Sharing require run-time work distribution and Scheduling

1	!\$OMP PARALLEL DO	#pragma omp for
2	do i=1,n	for(i=0;i<n;i++){
3	a(i)=b(i)+c(i)	a[i]=b[i]+c[i];
	enddo	}
5	!\$OMP END PARALLEL DO	

Line 1 Team of threads formed (parallel region).

Line 2-4 Loop iterations are split among threads.


Line 5 (Optional) end of parallel loop (implied barrier at enddo, }).

- Each loop iteration must be independent of other iterations.

Replicated and Work Share Constructs


- **Replicated:** Work blocks are executed by all threads.
- **Work Sharing:** Work is divided among threads.

```
PARALLEL
{code}
END PARALLEL
```




Replicated

```
PARALLEL DO
do I = 1,N*4
{code}
end do
END PARALLEL DO
```



Work Sharing

```
PARALLEL
{code1}
DO
do I = 1,N*4
{code2}
end do
{code3}
END PARALLEL
```



Combined

OpenMP Work-Sharing Scheduling

Clause Syntax: `SCHEDULE(schedule-type [,chunk-size])`

Schedule Type

– STATIC

- Threads receive chunks of iterations in thread order, round-robin
- Good if every iteration contains same amount of work
- May help keep parts of an array in a particular processor's cache— good between parallel do/for's.

– DYNAMIC

- Thread receives chunks as thread becomes available for more work
- Default chunk size is 1
- Good for load-balancing

OpenMP Work-Sharing Scheduling

– GUIDED

- Thread receives chunks as the thread becomes available for work
- Chunk size decreases exponentially, until it reaches the chunk size specified (default is 1)
- Balances load and reduces number of requests for more work

– RUNTIME

- Schedule is determined at run-time by `OMP_SCHEDULE`
- Useful for experimentation

OpenMP Work-Sharing Scheduling

For example, loop with 100 iterations and four threads

- SCHEDULE(STATIC)

Thread	0	1	2	3
Iteration	1-25	26-50	51-75	76-100

- SCHEDULE(DYNAMIC, 15) (*one possible outcome*)

Thread	0	1	3	2	1	3	2
Iteration	1-15	16-30	31-45	46-60	61-75	76-90	90-100

- SCHEDULE(GUIDED, 8) (*one possible outcome*)

Thread	0	1	2	3	3	2	3	1
Iteration	1-25	26-44	45-58	59-69	70-77	78-85	86-93	93-100

OpenMP Work-Sharing -- Sections

- **SECTIONS**

- Blocks of code are split among threads - task parallel style
- A thread might execute more than one block or no blocks
- Implied barrier

!\$OMP SECTIONS

```
!$OMP SECTION  
    CALL TASK1  
!$OMP SECTION  
    CALL TASK2  
!$OMP SECTION  
    CALL TASK3
```

!\$OMP END SECTIONS

```
#pragma omp sections  
{  
#pragma omp section  
    { TASK1( ); }  
#pragma omp section  
    { TASK2 ( ); }  
#pragma omp section  
    { TASK3 ( ); }  
}
```

OpenMP Work-Sharing -- Single

- **SINGLE**
 - Block of code is executed by a single thread
 - Implied barrier

!\$OMP single

**glob_count = glob_count + 1
print *, glob_count**

!\$OMP end single

#pragma single

**{
glob_count++;
printf(“%d\n”, glob_count);**

}

OpenMP Clauses -- Scoping

- Data scoping
 - PRIVATE
 - Each thread has its own copy of the specified variable
 - Variables are undefined after work sharing region
 - SHARED
 - Threads share a single copy of the specified variable
 - DEFAULT
 - A default of PRIVATE, SHARED or NONE can be specified
 - Note that loop counters of work sharing constructs are always PRIVATE by default; everything else is SHARED by default

OpenMP Data Scoping

- Data scoping (continued)
 - REDUCTION
 - Each thread has its own copy of the specified variable
 - Can appear only in reduction operation
 - All copies are "reduced" back into the original master thread's variable
 - FIRSTPRIVATE
 - Like PRIVATE, but copies are initialized using value from master thread's copy
 - LASTPRIVATE
 - Like PRIVATE, but final value is copied out to master thread's copy
 - For DO: last iteration; SECTIONS: last section

OpenMP Data Scoping

SHARED - Variable is shared (seen) by all processors.

PRIVATE - Each thread has a private instance of the variable.

Defaults: All DO LOOP indices are private, all other variables are shared. (OMP DO/FOR loops have private indices.)

```
!$omp parallel do shared(a), &  
                    private(t1,t2)
```

```
  do i = 1,100  
    t1 = f(i); t2 = g(i)  
    a(i) = sqrt(t1**2+t2**2)  
  end do
```

```
#pragma parallel for shared(a), \  
                    private(t1,t2)
```

```
  for(i=0; i<100; i++){  
    t1 = f[i]; t2 = g[i];  
    a[i] = sqrt( (t1*t1 + t2*t2);  
  }
```

All threads have access to the same storage areas for A, but each loop has its own private copy of the loop index, i, t1, and t2.

OpenMP Data Scoping

```
SUM = 0
!$OMP PARALLEL DO REDUCTION(+:SUM)
  DO I = 1, 100
    SUM = SUM + A(I)
  END DO
  ! Each thread's copy of SUM is added
  ! to original SUM at end of loop

!$OMP PARALLEL DO LASTPRIVATE(TEMP)
  DO I = 1, 100
    TEMP = F(I)
  END DO
  PRINT *, 'F(100) == ', TEMP
  ! TEMP is equal to F(100) at end of loop
```

OpenMP Work-Sharing Directives

- NOWAIT clause
 - Normally threads encounter a barrier synchronization at end of work-sharing construct.
 - Specifies that threads completing assigned work can proceed.
 - NOWAIT on END DO/END SECTIONS/END SINGLE line for FORTRAN. Include as clause in C/C++.